

INNOFarms - Comprehensive Code Documentation

Table of Contents

1. [High-Level Design \(HLD\)](#)
 2. [Low-Level Design \(LLD\)](#)
 3. [Module Documentation](#)
 - [Utils Module](#)
 - [Scrapers Module](#)
 - [ETL Module](#)
 - [Analytics Module](#)
 - [Financial Module](#)
 - [Visualization Module](#)
-


1. High-Level Design (HLD)


1.1 System Overview


The INNOFarms system is designed as a **data pipeline architecture** that follows the Extract-Transform-Load (ETL) pattern with additional analysis layers.

Data Flow Diagram









1.2 Design Principles

1.2.1 Separation of Concerns

Each module handles one specific responsibility:

- Scrapers: Data collection only
- ETL: Data cleaning only
- Analytics: Analysis only
- Financial: Business logic only

1.2.2 Dependency Injection

- Configuration is injected via `.env` and YAML files

- Database path is configurable
- Logging is centralized

1.2.3 Modularity

- New data sources: Create new scraper class
- New commodities: Edit YAML config
- New analysis: Add new method to analyzer


1.2.4 Error Handling


- Each operation logs errors
 - Failed records don't crash pipeline
 - Graceful degradation with sample data
-

2. Low-Level Design (LLD)


2.1 Class Diagrams

2.1.1 Scraper Module





2.1.2 ETL Pipeline



```
| - _create_schema(conn) |  
| - _get_commodity_map(conn, commodities) → Dict |  
| - _get_source_map(conn, sources) → Dict |  
| - _guess_category(commodity) → str |  
| - _create_aggregates(conn) |
```

2.1.3 Analytics Module

```
| TrendAnalyzer |  
|  
| - db_path: str |  
| - results: Dict[str, Any] |  
|  
| + get_price_series(commodity, days) → DataFrame |  
| + calculate_moving_averages(df, windows) → Dict |  
| + detect_trend(df, short, long) → Dict |  
| + calculate_volatility(df) → Dict |  
| + analyze_commodity(commodity, days) → Dict |  
| + analyze_all(commodities, days) → Dict |  
| + generate_report(output_file) → str |
```

```
| ForecastEngine |  
|  
| - db_path: str |  
| - forecasters: Dict[str, Forecaster] |  
| - results: Dict[str, Any] |  
|  
| + get_price_series(commodity, days) → Series |  
| + evaluate_model(series, forecaster) → Dict |
```

```

| + forecast_commodity(commodity, days) → Dict      |
| + forecast_all(commodities, days) → Dict          |
| + generate_report(output_file) → str            |

```

△

| uses

| MAForecaster |

| ARIMAForecaster |


| + window: int |


| + order: tuple |

| + forecast() |

| + forecast() |

2.2 Database Schema (ER Diagram)






| aggregated by




└──+ daily_aggregates ─────────────────────────




2.3 Sequence Diagrams

2.3.1 Data Scraping Sequence



2.3.2 ETL Pipeline Sequence





3. Module Documentation

3.1 Utils Module

3.1.1 logger.py - Centralized Logging

```
"""
```

Centralized Logger Configuration

```
=====
```

Provides a unified logging configuration for all modules.

Import this module to get a configured logger.

Usage:

```
from src.utils.logger import get_logger
logger = get_logger(__name__)
logger.info("Message here")
"""
```

```
import os
import logging
from datetime import datetime
from pathlib import Path
from typing import Optional
```

Line-by-line explanation:

1. `""" ... """` - Module docstring explaining purpose and usage
2. `import os` - Operating system interface for file paths
3. `import logging` - Python's built-in logging module
4. `from datetime import datetime` - For timestamp generation
5. `from pathlib import Path` - Modern path handling
6. `from typing import Optional` - Type hints for optional parameters

```
# =====
# CONFIGURATION
# =====

LOG_DIR = Path(__file__).parent.parent.parent / "logs"
LOG_FORMAT = "%(asctime)s - %(name)s - %(levelname)s - %(message)s"
DATE_FORMAT = "%Y-%m-%d %H:%M:%S"
DEFAULT_LEVEL = logging.INFO
```

Why these constants?

- `LOG_DIR`: Places logs in project root's `logs/` folder
 - `Path(__file__)` = current file path
 - `.parent.parent.parent` = go up 3 directories (`utils → src → project root`)
- `LOG_FORMAT`: Standard logging format showing:
 - `%(asctime)s` - Timestamp
 - `%(name)s` - Logger name (usually module name)
 - `%(levelname)s` - DEBUG/INFO/WARNING/ERROR
 - `%(message)s` - Actual log message
- `DATE_FORMAT`: ISO-like format for readability
- `DEFAULT_LEVEL`: INFO level shows important events, not debug noise

```
def get_logger(
    name: str,
    level: int = DEFAULT_LEVEL,
    log_to_file: bool = True,
```

```
    log_file: Optional[str] = None
) -> logging.Logger:
    """
Get a configured logger instance.

Args:
    name: Logger name (typically __name__)
    level: Logging level (default: INFO)
    log_to_file: Whether to also log to file
    log_file: Custom log file path (optional)
```

Returns:

Configured logger instance

"""

Function signature explanation:

- `name: str` - Logger identifier, use `__name__` to get module name automatically
- `level: int` - Control verbosity (DEBUG=10, INFO=20, WARNING=30, ERROR=40)
- `log_to_file: bool` - Enable file logging for debugging
- `log_file: Optional[str]` - Custom path, or auto-generate from name
- `-> logging.Logger` - Returns configured logger object

```
logger = logging.getLogger(name)

# Only configure if not already configured
if logger.handlers:
    return logger
```

Why check handlers?

- `getLogger(name)` returns same object for same name (singleton pattern)
- If already configured (has handlers), don't add duplicate handlers
- Prevents duplicate log messages when importing module multiple times

```
logger.setLevel(level)

# Console handler
console_handler = logging.StreamHandler()
console_handler.setLevel(level)
console_handler.setFormatter(logging.Formatter(LOG_FORMAT, DATE_FORMAT))
logger.addHandler(console_handler)
```

Console handler setup:

1. `StreamHandler()` - Outputs to stdout (terminal)
2. `setLevel()` - Filter messages below this level
3. `setFormatter()` - Apply our format string
4. `addHandler()` - Attach to logger

```
# File handler (optional)
if log_to_file:
    LOG_DIR.mkdir(parents=True, exist_ok=True)

if log_file is None:
    # Use module name as log file
    safe_name = name.replace(".", "_").replace("/", "_")
    log_file = LOG_DIR / f"{safe_name} {datetime.now().strftime('%Y%m%d')}.log"
else:
    log_file = Path(log_file)

file_handler = logging.FileHandler(log_file)
file_handler.setLevel(level)
file_handler.setFormatter(logging.Formatter(LOG_FORMAT, DATE_FORMAT))
logger.addHandler(file_handler)

return logger
```

File handler setup:

1. `mkdir(parents=True, exist_ok=True)` - Create logs directory if missing
 2. `safe_name` - Replace dots and slashes with underscores for valid filename
 3. `datetime.now().strftime('%Y%m%d')` - Daily log rotation by date
 4. `FileHandler` - Writes to disk for future debugging
-

3.1.2 config.py - Environment Configuration

```
"""  
Configuration loader with environment variable support.  
"""
```

```
import os  
from pathlib import Path  
from typing import Dict, Any, Optional  
  
# Try to load python-dotenv  
try:  
    from dotenv import load_dotenv  
    load_dotenv()  
except ImportError:  
    pass
```

Why try/except for dotenv?

- `python-dotenv` loads `.env` file into environment variables
- If not installed, fail silently (doesn't crash the app)
- `load_dotenv()` reads `.env` file and sets `os.environ` values

```
def get_env(key: str, default: Optional[str] = None) -> Optional[str]:  
    """Get environment variable with optional default."""  
    return os.environ.get(key, default)
```

```

def get_env_int(key: str, default: int = 0) -> int:
    """Get environment variable as integer."""
    value = os.environ.get(key)
    return int(value) if value else default

def get_env_float(key: str, default: float = 0.0) -> float:
    """Get environment variable as float."""
    value = os.environ.get(key)
    return float(value) if value else default

def get_env_bool(key: str, default: bool = False) -> bool:
    """Get environment variable as boolean."""
    value = os.environ.get(key, "").lower()
    if value in ("true", "1", "yes"):
        return True
    if value in ("false", "0", "no"):
        return False
    return default

```

Helper functions explained:

- `get_env()` - Basic string retrieval with default
- `get_env_int()` - Parse string to int (e.g., “5” → 5)
- `get_env_float()` - Parse string to float (e.g., “3.14” → 3.14)
- `get_env_bool()` - Handle various boolean representations

Why these helpers?

- Environment variables are always strings
- Type conversion with error handling
- Consistent default value pattern

```

class Config:
    """Application configuration from environment variables."""

    # API Keys
    USDA_API_KEY: str = get_env("USDA_API_KEY", "")

    # Database
    DATABASE_PATH: str = get_env("DATABASE_PATH", "database/market_data.d
        b")

    # Logging
    LOG_LEVEL: str = get_env("LOG_LEVEL", "INFO")

    # Scraper settings
    RATE_LIMIT_MIN: int = get_env_int("SCRAPER_RATE_LIMIT_MIN", 2)
    RATE_LIMIT_MAX: int = get_env_int("SCRAPER_RATE_LIMIT_MAX", 5)
    MAX_RETRIES: int = get_env_int("SCRAPER_MAX_RETRIES", 3)

    # Paths
    PROJECT_ROOT: Path = Path(__file__).parent.parent.parent
    DATA_DIR: Path = PROJECT_ROOT / "data" / "raw"
    REPORTS_DIR: Path = PROJECT_ROOT / "reports"
    CONFIG_DIR: Path = PROJECT_ROOT / "config"

config = Config()

```

Why a Config class?

- **Single source of truth** - All settings in one place
- **Type hints** - IDE autocomplete and error checking
- **Default values** - Works without .env file
- **Singleton pattern** - `config = Config()` creates single instance

3.2 Scrapers Module

3.2.1 usda_scraper.py - USDA NASS API Scraper

```
# =====
# CONSTANTS
# =====
USDA_API_URL = "https://quickstats.nass.usda.gov/api/api_GET/"

# Commodity configurations - easily extendable
COMMODITIES = {
    "WHEAT": {
        "commodity_desc": "WHEAT",
        "statisticcat_desc": "PRICE RECEIVED",
        "unit_desc": "$ / BU"
    },
    "CORN": {
        "commodity_desc": "CORN",
        "statisticcat_desc": "PRICE RECEIVED",
        "unit_desc": "$ / BU"
    },
    # ... more commodities
}
```

Constants design:

- `USDA_API_URL` - API endpoint (could move to config)
- `COMMODITIES` - Dictionary for each commodity's API parameters
 - `commodity_desc` - USDA commodity name
 - `statisticcat_desc` - What statistic to fetch (price, production, etc.)
 - `unit_desc` - Expected unit for price

Why dictionary format?

- Easy to add new commodities

- Self-documenting (field names explain purpose)
- Can be overridden per-call

```
class USDANASSScraper:
    """
    Scraper for USDA NASS QuickStats API.
    """

    def __init__(self, api_key: str = None):
        """
        Initialize the USDA scraper.

        Args:
            api_key: USDA API key (defaults to env variable)
        """

        self.api_key = api_key or config.USDA_API_KEY
        self.base_url = USDA_API_URL
        self.data: Dict[str, List[Dict]] = {}
        self.rate_limit_min = config.RATE_LIMIT_MIN
        self.rate_limit_max = config.RATE_LIMIT_MAX

        if not self.api_key:
            logger.warning("USDA API key not set")
```

Constructor explained:

- `api_key or config.USDA_API_KEY` - Use provided key OR fall back to env
- `self.data` - Stores scraped data (key: commodity name, value: records list)
- Rate limits from config - respect API terms of service
- Warning if no API key - helps debug issues

```
def _make_request(self, params: Dict) -> Optional[Dict]:
    """
    Make API request with error handling.
```

```

"""
params['key'] = self.api_key
params['format'] = 'JSON'

try:
    response = requests.get(
        self.base_url,
        params=params,
        timeout=30
    )
    response.raise_for_status()
    return response.json()

except requests.exceptions.RequestException as e:
    logger.error(f"API request failed: {e}")
    return None

```

HTTP request pattern:

1. Add authentication (`key`) and format to params
2. `requests.get()` - Make HTTP GET request
3. `timeout=30` - Don't hang forever
4. `raise_for_status()` - Raise exception for 4xx/5xx errors
5. Return `None` on error (caller handles gracefully)

```

def _rate_limit(self):
    """Apply rate limiting delay."""
    delay = random.uniform(self.rate_limit_min, self.rate_limit_max)
    time.sleep(delay)

```

Why rate limiting?

- Prevents overwhelming the API server
- Avoids getting IP blocked
- `random.uniform()` - Random delay looks more human-like
- Typically 2-5 seconds between requests

```
def scrape_commodity(
    self,
    commodity: str,
    years: int = 10,
    commodity_config: Dict = None
) -> List[Dict]:
    """
    Scrape data for a single commodity.
    """

    if commodity_config is None:
        commodity_config = COMMODITIES.get(commodity.upper(), {})

    if not commodity_config:
        logger.warning(f"No configuration for commodity: {commodity}")
        return []

    current_year = datetime.now().year
    start_year = current_year - years

    params = {
        "source_desc": "SURVEY",
        "sector_desc": "CROPS",
        "commodity_desc": commodity_config.get("commodity_desc", commodity),
        "statisticcat_desc": commodity_config.get("statisticcat_desc", "PRICE RECEIVED"),
        "year_GE": str(start_year),
        "year_LE": str(current_year),
        "freq_desc": "ANNUAL"
    }

    logger.info(f"Fetching {commodity} data ({start_year}-{current_year})...")

    response = self._make_request(params)
```

Query building:

- USDA API uses filter parameters:
 - `source_desc="SURVEY"` - Survey data (not census)
 - `sector_desc="CROPS"` - Crop sector (not livestock)
 - `year_GE` - Year greater-than-or-equal (start)
 - `year_LE` - Year less-than-or-equal (end)
- Uses commodity config for commodity-specific parameters

```
if not response or 'data' not in response:
    logger.warning(f"No data returned for {commodity}")
    return []

records = []
for item in response['data']:
    try:
        value = item.get('Value', "").replace(',', '')
        if value and value != '(D)' and value != '(NA)':
            record = {
                'commodity': commodity.title(),
                'date': f"{item.get('year', datetime.now().year)}-01-01",
                'price': float(value),
                'unit': item.get('unit_desc', 'USD'),
                'source': 'USDA NASS',
                'state': item.get('state_name', 'US TOTAL'),
                'market': 'National Average',
                'frequency': item.get('freq_desc', 'ANNUAL')
            }
            records.append(record)
    except (ValueError, KeyError) as e:
        logger.debug(f"Skipping invalid record: {e}")
        continue

logger.info(f"Scraped {len(records)} records for {commodity}")
self._rate_limit()
```

```
    return records
```

Record parsing:

- `value.replace(',', '')` - Remove thousands separator (1,234 → 1234)
- `value != '(D)'` - Skip suppressed data
- `value != '(NA)'` - Skip not available
- `try/except` - Skip malformed records without crashing
- Normalize field names (commodity, date, price, etc.)

```
def add_commodity(self, name: str, config_dict: Dict) -> None:
```

```
    """
```

Add a new commodity configuration for scraping.

```
    """
```

```
    COMMODITIES[name.upper()] = config_dict
    logger.info(f"Added commodity: {name}")
```

Why add_commodity method?

- Allows runtime addition without editing source
- Extensibility for users
- Logs the addition for debugging

3.3 ETL Module

3.3.1 etl_pipeline.py - Extract Transform Load

This is the core data processing module. Let me explain the key methods:

```
def extract(self, file_pattern: str = "*.json") -> List[Dict]:
```

```
    """
```

Extract data from raw JSON files.

```
    """
```

```

logger.info(f"Extracting data from {self.input_dir}")

all_records = []
json_files = list(self.input_dir.glob(file_pattern))

logger.info(f"Found {len(json_files)} JSON files")

for json_file in json_files:
    try:
        with open(json_file) as f:
            data = json.load(f)

        # Extract source
        source = data.get('source', 'Unknown')

        # Extract records from commodities dict
        commodities_data = data.get('commodities', {})
        for commodity_name, records in commodities_data.items():
            if isinstance(records, list):
                for record in records:
                    record['source'] = source
                    record['commodity'] = commodity_name
                    all_records.append(record)

        self.stats['files_processed'] += 1

    except Exception as e:
        logger.error(f"Error extracting {json_file}: {e}")

    self.stats['records_extracted'] = len(all_records)

```

Extract phase explained:

1. `glob(file_pattern)` - Find all matching files

2. Loop through each JSON file
3. Parse nested structure (commodities → array of records)
4. Flatten into single list of records
5. Add source/commodity to each record
6. Count statistics

```
def transform(self, records: List[Dict]) -> pd.DataFrame:
    """
    Transform extracted records.
    """

    if not records:
        return pd.DataFrame()

    df = pd.DataFrame(records)
    initial_count = len(df)

    # Normalize column names
    df.columns = df.columns.str.lower().str.strip()
```

Why normalize columns?

- Different sources use different cases (Price vs price vs PRICE)
- `str.lower()` - Standardize to lowercase
- `str.strip()` - Remove whitespace

```
# Parse dates
df['date'] = pd.to_datetime(df['date'], errors='coerce')
df = df.dropna(subset=['date'])
```

Date handling:

- `pd.to_datetime()` - Parse various date formats
- `errors='coerce'` - Invalid dates become NaT (null)
- `dropna(subset=['date'])` - Remove rows with invalid dates

```

# Remove duplicates
before_dedupe = len(df)
df = df.drop_duplicates(
    subset=['commodity', 'date', 'source', 'market'],
    keep='first'
)
self.stats['duplicates_removed'] = before_dedupe - len(df)

```

Deduplication:

- Same commodity+date+source+market = duplicate
- `keep='first'` - Keep first occurrence
- Track duplicates removed for statistics

```

# Detect anomalies (prices > 3 std from mean)
for commodity in df['commodity'].unique():
    mask = df['commodity'] == commodity
    prices = df.loc[mask, 'price']
    if len(prices) > 10:
        mean, std = prices.mean(), prices.std()
        if std > 0:
            z_scores = abs((prices - mean) / std)
            anomaly_mask = z_scores > 3
            anomaly_count = anomaly_mask.sum()
            if anomaly_count > 0:
                self.stats['anomalies_detected'] += anomaly_count

```

Anomaly detection:

- Z-score = $(\text{value} - \text{mean}) / \text{standard deviation}$
- $|Z| > 3$ means value is >3 standard deviations from mean
- 99.7% of normal data falls within 3 std
- Flag (but don't remove) anomalies

```

# Create record hash for deduplication
df['record_hash'] = df.apply(
    lambda r: hashlib.md5(
        f"{r['commodity']}{r['date']}{r['source']}{r.get('market', '')}".encode()
    ).hexdigest()[:16],
    axis=1
)

```

Why record hash?

- Unique identifier for each record
 - Used as database primary key
 - MD5 produces consistent hash for same input
 - [:16] - First 16 chars is enough for uniqueness
-

3.4 Analytics Module

3.4.1 trend_analyzer.py

```

def calculate_moving_averages(
    self,
    df: pd.DataFrame,
    windows: List[int] = [7, 15, 30]
) -> Dict[str, float]:
    """
    Calculate moving averages for different windows.
    """
    result = {}

    for window in windows:
        if len(df) >= window:
            ma = df['price'].rolling(window=window).mean().iloc[-1]
            result[window] = ma

```

```

        result[f'ma_{window}_day'] = round(ma, 2)
    else:
        result[f'ma_{window}_day'] = None

    return result

```

Moving averages explained:

- `rolling(window=7)` - Create 7-day sliding window
- `.mean()` - Average of each window
- `.iloc[-1]` - Get latest (most recent) value
- Different windows show different trends:
 - 7-day: Short-term, responsive to recent changes
 - 15-day: Medium-term
 - 30-day: Long-term, smooth out noise

Mathematical formula:

$$MA_n = (P_1 + P_2 + \dots + P_n) / n$$

```

def detect_trend(
    self,
    df: pd.DataFrame,
    short_window: int = 7,
    long_window: int = 30
) -> Dict[str, Any]:
    """
    Detect price trend direction.
    """

    if len(df) < long_window:
        return {'trend': 'Unknown', 'strength': 0}

    short_ma = df['price'].rolling(window=short_window).mean().iloc[-1]
    long_ma = df['price'].rolling(window=long_window).mean().iloc[-1]

```

```

if long_ma == 0:
    return {'trend': 'Unknown', 'strength': 0}

diff_pct = ((short_ma - long_ma) / long_ma) * 100

if diff_pct > 2:
    trend = 'Upward'
elif diff_pct < -2:
    trend = 'Downward'
else:
    trend = 'Stable'

return {
    'trend': trend,
    'strength': round(abs(diff_pct), 2)
}

```

Trend detection logic:

1. Calculate short-term MA (responsive to recent changes)
2. Calculate long-term MA (baseline)
3. Compare: $(\text{short} - \text{long}) / \text{long} * 100$
4. If short > long by >2%: Upward trend
5. If short < long by >2%: Downward trend
6. Otherwise: Stable

Why 2% threshold?

- Small differences could be noise
- 2% is significant enough to indicate real trend
- Adjustable based on commodity volatility

```
def calculate_volatility(self, df: pd.DataFrame) -> Dict[str, Any]:
```

```
    """
```

Calculate price volatility using coefficient of variation.

```

"""
if df.empty or len(df) < 10:
    return {'volatility_pct': 0, 'level': 'Unknown'}

mean_price = df['price'].mean()
std_price = df['price'].std()

if mean_price == 0:
    return {'volatility_pct': 0, 'level': 'Unknown'}

cv = (std_price / mean_price) * 100

if cv < 10:
    level = 'Low'
elif cv < 20:
    level = 'Medium'
else:
    level = 'High'

return {
    'volatility_pct': round(cv, 2),
    'level': level
}

```

Coefficient of Variation (CV):

$$CV = (\text{Standard Deviation} / \text{Mean}) \times 100\%$$

Why CV instead of just standard deviation?

- CV is normalized by mean
- Allows comparison between different price scales
- \$100 commodity with \$10 std vs \$1000 commodity with \$10 std
- Same std, but very different volatility!

3.4.2 forecasters.py

```
class MovingAverageForecaster:  
    """Simple Moving Average forecaster."""  
  
    def __init__(self, window: int = 7):  
        self.window = window  
        self.name = "Moving Average"  
  
    def forecast(self, series: pd.Series, periods: int = 7) -> List[float]:  
        """Generate forecast using moving average trend."""  
        if len(series) < self.window:  
            return [series.mean()] * periods  
  
        ma = series.tail(self.window).mean()  
        trend = (series.iloc[-1] - series.iloc[-self.window]) / self.window  
  
        forecasts = []  
        for i in range(1, periods + 1):  
            forecasts.append(round(ma + trend * i, 2))  
  
        return forecasts
```

MA Forecasting logic:

1. Calculate recent MA (last 7 days)
2. Calculate trend: $(\text{latest price} - \text{price 7 days ago}) / 7$
3. Project forward: $\text{MA} + \text{trend} * \text{days_ahead}$

Limitations:

- Assumes linear trend continues
- Doesn't capture seasonality
- Simple but fast baseline

```

class ARIMAForecaster:
    """ARIMA time series forecaster."""

    def __init__(self, order: Tuple[int, int, int] = (5, 1, 0)):
        self.order = order
        self.name = "ARIMA"

    def forecast(self, series: pd.Series, periods: int = 7) -> List[float]:
        """Generate ARIMA forecast."""
        try:
            from statsmodels.tsa.arima.model import ARIMA

            if len(series) < 30:
                return [series.mean()] * periods

            model = ARIMA(series, order=self.order)
            fitted = model.fit()
            forecasts = fitted.forecast(steps=periods)

            return [round(f, 2) for f in forecasts.tolist()]

        except Exception as e:
            logger.warning(f"ARIMA failed: {e}, using fallback")
            return [round(series.mean(), 2)] * periods

```

ARIMA explained:

- **AR(p)**: AutoRegressive - uses past values
- **I(d)**: Integrated - differencing for stationarity
- **MA(q)**: Moving Average - uses past errors

Order (5, 1, 0) means:

- p=5: Use last 5 values for autoregression
- d=1: First difference (makes non-stationary data stationary)
- q=0: No moving average of errors

Why (5, 1, 0)?

- Simple model that often works
- 1 differencing handles linear trends
- 5 lags capture weekly patterns

```
def calculate_mape(actual: np.ndarray, predicted: np.ndarray) -> float:  
    """Calculate Mean Absolute Percentage Error."""  
    actual = np.array(actual)  
    predicted = np.array(predicted)  
  
    mask = actual != 0  
    if not mask.any():  
        return 0.0  
  
    mape = np.mean(np.abs((actual[mask] - predicted[mask]) / actual[mask])) * 10  
          0  
    return round(mape, 2)
```

MAPE formula:

$$\text{MAPE} = \frac{1}{n} \times \sum |actual - predicted| / actual \times 100\%$$

Why MAPE?

- Percentage-based (scale-independent)
 - Easy to interpret (5% error means 5% off on average)
 - Mask zeros to avoid division by zero
-

3.5 Financial Module

3.5.1 cost_calculator.py

```
class CostCalculator:  
    """  
  
    Calculates production costs per hectare for commodities.  
    """  
  
    def __init__(self, config_path: str = "config/production_costs.yaml"):  
        self.config = self._load_config(config_path)  
  
    def _load_config(self, path: str) -> Dict:  
        """Load cost configuration from YAML."""  
        try:  
            with open(path) as f:  
                return yaml.safe_load(f)  
        except:  
            return self._get_default_costs()  
  
    def get_total_cost(self, commodity: str) -> float:  
        """  
        Get total production cost per hectare.  
        """  
        costs = self.config.get('costs', {}).get(commodity, {})  
  
        return sum([  
            costs.get('seeds', 0),  
            costs.get('fertilizers', 0),  
            costs.get('pesticides', 0),  
            costs.get('labor', 0),  
            costs.get('irrigation', 0),  
            costs.get('machinery', 0),
```

```
    costs.get('other', 0)
])
```

Cost categories:

1. **Seeds**: Initial planting material
 2. **Fertilizers**: NPK, micronutrients
 3. **Pesticides**: Herbicides, insecticides, fungicides
 4. **Labor**: Planting, weeding, harvesting
 5. **Irrigation**: Water, pumping costs
 6. **Machinery**: Tractors, equipment rental
 7. **Other**: Transport, storage, misc
-

3.5.2 profit_calculator.py

```
def calculate_profit(
    self,
    commodity: str,
    price_per_unit: float = None,
    yield_per_hectare: float = None
) -> Dict[str, float]:
```

```
    """
```

Calculate profit for a commodity.

Returns:

Dict with revenue, cost, profit, roi, margin

```
    """
```

Get current price if not provided

if price_per_unit is None:

```
        price_per_unit = self._get_current_price(commodity)
```

Get expected yield if not provided

if yield_per_hectare is None:

```
        yield_per_hectare = self._get_expected_yield(commodity)
```

```

# Calculate
total_cost = self.cost_calc.get_total_cost(commodity)
revenue = price_per_unit * yield_per_hectare
profit = revenue - total_cost

roi = (profit / total_cost) * 100 if total_cost > 0 else 0
margin = (profit / revenue) * 100 if revenue > 0 else 0

return {
    'commodity': commodity,
    'price_per_unit': price_per_unit,
    'yield_per_hectare': yield_per_hectare,
    'revenue': round(revenue, 2),
    'total_cost': round(total_cost, 2),
    'profit': round(profit, 2),
    'roi_percent': round(roi, 2),
    'profit_margin_percent': round(margin, 2)
}

```

Financial formulas:

Revenue = Price × Yield
 Profit = Revenue - Cost
 ROI = (Profit / Cost) × 100%
 Margin = (Profit / Revenue) × 100%

Why both ROI and Margin?

- **ROI:** Return relative to investment (cost)
 - Answers: “How much do I get back per rupee invested?”
 - Higher is better for capital efficiency
- **Margin:** Profit relative to sales (revenue)

- Answers: “What % of revenue is profit?”
 - Important for pricing decisions
-

3.5.3 opportunity_scorer.py

```
def calculate_score(self, commodity: str) -> Dict[str, Any]:  
    """  
    Calculate opportunity score for a commodity.  
  
    Score = weighted sum of:  
    - ROI score (35%)  
    - Trend score (25%)  
    - Stability score (20%)  
    - Demand score (20%)  
    """  
  
    # Get component scores  
    roi_score = self._get_roi_score(commodity)  
    trend_score = self._get_trend_score(commodity)  
    stability_score = self._get_stability_score(commodity)  
    demand_score = self._get_demand_score(commodity)  
  
    # Weighted sum  
    total_score = (  
        roi_score * 0.35 +  
        trend_score * 0.25 +  
        stability_score * 0.20 +  
        demand_score * 0.20  
    )  
  
    # Determine recommendation  
    if total_score >= 70:  
        recommendation = "Strong Buy"  
    elif total_score >= 50:
```

```

recommendation = "Buy"
elif total_score >= 30:
    recommendation = "Hold"
else:
    recommendation = "Avoid"

return {
    'commodity': commodity,
    'score': round(total_score, 1),
    'recommendation': recommendation,
    'components': {
        'roi': roi_score,
        'trend': trend_score,
        'stability': stability_score,
        'demand': demand_score
    }
}

```

Scoring methodology:

1. **ROI Score (35%)** - Higher ROI = higher score
2. **Trend Score (25%)** - Upward trend = higher score
3. **Stability Score (20%)** - Lower volatility = higher score
4. **Demand Score (20%)** - Market demand indicators

Why these weights?

- ROI is most important (direct profit indicator)
 - Trend shows future potential
 - Stability reduces risk
 - Demand ensures market exists
-

3.6 Visualization Module

3.6.1 visualizer.py

```
def plot_price_trends(
    self,
    commodity: str,
    days: int = 180
) -> Optional[str]:
    """
    Plot price trends with moving averages.
    """

    # ... query data ...

    df['ma_7'] = df['price'].rolling(7).mean()
    df['ma_30'] = df['price'].rolling(30).mean()

    fig, ax = plt.subplots(figsize=(12, 6))

    ax.plot(df['price_date'], df['price'],
            label='Daily Price', alpha=0.7, linewidth=1)
    ax.plot(df['price_date'], df['ma_7'],
            label='7-Day MA', linewidth=2)
    ax.plot(df['price_date'], df['ma_30'],
            label='30-Day MA', linewidth=2)

    ax.fill_between(df['price_date'], df['price'], alpha=0.1)

    ax.set_title(f'{commodity} Price Trends ({days} Days)', fontweight='bold')
    ax.set_xlabel('Date')
    ax.set_ylabel('Price')
    ax.legend(loc='upper left')
    ax.grid(True, alpha=0.3)
```

Matplotlib components:

- `plt.subplots()` - Create figure and axes
 - `ax.plot()` - Line plot
 - `ax.fill_between()` - Shaded area under curve
 - `ax.set_title()` - Chart title
 - `ax.legend()` - Show label legend
 - `ax.grid()` - Add gridlines
-

Summary

This documentation covered:

1. **High-Level Design** - System architecture, data flow, design principles
2. **Low-Level Design** - Class diagrams, database schema, sequence diagrams
3. **Module Documentation** - Line-by-line code explanation for each module

Key design patterns used:

- **Strategy Pattern** - Multiple scrapers with same interface
- **Pipeline Pattern** - ETL with distinct phases
- **Singleton Pattern** - Config and logger instances
- **Factory Pattern** - Logger creation

Key algorithms:

- **Moving Average** - Trend smoothing
 - **ARIMA** - Time series forecasting
 - **Coefficient of Variation** - Volatility measurement
 - **Weighted Scoring** - Multi-criteria decision making
-

Last updated: February 1, 2026